

Logistic Regression

(part II)

To perform inference in Logistic Reg,
we must rely on asymptotic
results (i.e. using a large sample
approximation).

I Wald test for indiv coefficients.

Main Idea: When n is large, the MLE

$$\hat{\beta}_k \stackrel{\text{(approx)}}{\sim} N(\beta_k, \text{Var}(\hat{\beta}_k))$$

where $\text{Var}(\hat{\beta}_k)$ is defined as follows:

┌

Let G denote the Hessian matrix
of the loglikelihood:

$$G = \begin{bmatrix} \frac{\partial^2 \ell}{\partial \beta_0^2} & \frac{\partial^2 \ell}{\partial \beta_0 \partial \beta_1} & \dots & \frac{\partial^2 \ell}{\partial \beta_0 \partial \beta_{p-1}} \\ \vdots & \frac{\partial^2 \ell}{\partial \beta_1^2} & \ddots & \vdots \\ \frac{\partial^2 \ell}{\partial \beta_0 \partial \beta_{p-1}} & \frac{\partial^2 \ell}{\partial \beta_1 \partial \beta_{p-1}} & \dots & \frac{\partial^2 \ell}{\partial \beta_{p-1}^2} \end{bmatrix}$$

then the variance of the MLE is:

$$\text{Var}(\hat{\beta}) = -G^{-1} \Big|_{\beta = \hat{\beta}}$$

↳

$$\text{Var}(\hat{\beta}_k) = \left[-G^{-1} \Big|_{\beta = \hat{\beta}} \right]_{k+1, k+1}$$

where $\hat{\beta}$ is the MLE of β .

OK! So then

$$\frac{\hat{\beta}_k - \beta_k}{SE(\hat{\beta}_k)} \sim N(0, 1) \quad \&$$

a large sample test for β_k can be constructed as:

$$H_0: \beta_k = 0 \quad \text{vs} \quad H_1: \beta_k \neq 0 \quad \text{w/}$$

The test statistic:

$$Z = \frac{\hat{\beta}_k}{SE(\hat{\beta}_k)} \quad \& \quad \text{the decision rule,}$$

if $|Z| > z_{1-\alpha/2}^*$, reject H_0 .

II Deviance & Likelihood Ratio Tests for Reduced & Full Models.

If we want to compare two models w/
a different # of predictors, for ex:

$$H_0: \beta_r = \beta_{r+1} = \dots = \beta_{p-1} = 0$$

vs.

$$H_1: \text{at least one } \beta_j, j=r, \dots, p-1, \text{ not zero}$$



$$H_0: \text{Reduced Model} \Rightarrow \text{logit}(\hat{\pi}_i) = \beta_0 + \sum_{j=r}^{p-1} \beta_j x_{ji}$$

$$H_1: \text{Full Model} \Rightarrow \text{logit}(\hat{\pi}_i) = \beta_0 + \sum_{j=1}^{p-1} \beta_j x_{ji}$$

Our test is based on comparing the log-likelihoods
across H_0 & H_1 .

Deviance

Def: $\text{Deviance} = -2 \ln(\beta)$

Then a likelihood-based test Stat is defined as:

$$\Delta D = \text{Deviance (Reduced)} - \text{Deviance (Full)}$$

Decision Rule:

If $\Delta D > \chi^2_{df, 1-\alpha}$, reject H_0

Here $df = P_{\text{Full}} - P_{\text{Red}}$.

If we reject H_0 , we have evidence that the full model is a significantly better fit according to likelihood.

Recall, for logistic reg,

$$l = \sum_{i=1}^n \left(y_i \log \left(\frac{\pi_i}{1-\pi_i} \right) + \log(1-\pi_i) \right)$$

So deviance of a fitted model is:

$$\text{Dev}(\hat{\beta}) = -2 \sum_{i=1}^n \left(y_i \log \left(\frac{\hat{\pi}_i}{1-\hat{\pi}_i} \right) + \log(1-\hat{\pi}_i) \right)$$

III Deviance Residuals

OLS residuals don't apply to logistic reg.

Observed response: $y_i = \begin{cases} 1 \\ 0 \end{cases}$

predicted response: $\hat{\pi}_i = \hat{p}(y_i=1)$


If we want to have an analog of "residuals" there are 2 alternatives:

① "Pearson Residuals"

— similar to studentized resid in OLS

$$e_i = \frac{y_i - \hat{\pi}_i}{\sqrt{\hat{\pi}_i(1 - \hat{\pi}_i)}} = \frac{y_i - \hat{\pi}_i}{\sqrt{\hat{\pi}_i(1 - \hat{\pi}_i)}}$$

We would like to observe Pearson residuals scattered randomly & evenly across the x-axis:



↳ interpret like usual residual plots!!


② Deviance Residuals

Define the deviance residuals as:

$$d_i = \begin{cases} \sqrt{-2(y_i \log(\hat{\pi}_i / (1 - \hat{\pi}_i)) + \log(1 - \hat{\pi}_i))} & \text{if } y_i = 1, \\ -\sqrt{-2(y_i \log(\hat{\pi}_i / (1 - \hat{\pi}_i)) + \log(1 - \hat{\pi}_i))} & \text{if } y_i = 0. \end{cases}$$

The square of each deviance residual measures the contribution of each response to the deviance of the fitted model.

We can check the dev-resid scatter plot:



- Pseudo R^2 :

Because there is no OLS principle, the regular R^2 doesn't exist for logit reg. to "explain variance".

Some pseudo R^2 statistics have been invented to assess goodness-of-fit.

Efron:

$$\text{pseudo } R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{\pi}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2}$$

McFadden: pseudo $R^2 = 1 - \frac{l(\text{Full})}{l(\text{Null})}$

Note: AIC/BIC can still be used
for model selection w/ logistic reg.